

What are support (SVMs)?

Author



Eda Kavlakoglu

Business Development +
Partnerships
IBM Research

What are SVMs?



A support vector machine (SVM) is a algorithm that classifies data by find that maximizes the distance between space.

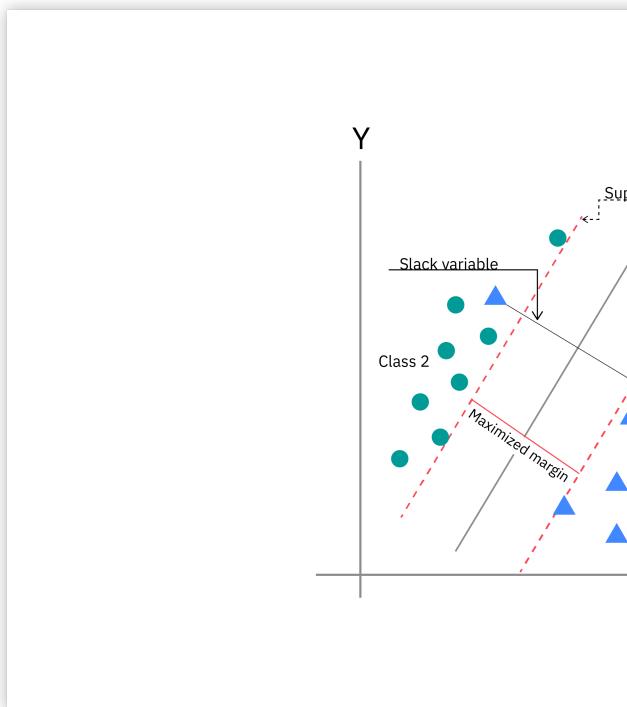
SVMs were developed in the 1990s by Vladimir Vapnik. He published this work in a paper titled "Support Vector Method for Function Estimation, Classification, Regression Estimation, and Signal Processing".

SVMs are commonly used within classification and regression analysis. Your privacy choices Optimal hyperplane that separates data points of opposite classes. The number of support vectors is minimal.



hyperplane is a line in a 2-D space or a plane
hyperplanes can be found to differentiate clas
enables the algorithm to find the best decisio
enables it to generalize well to new data and
lines that are adjacent to the optimal hyperpl
vectors run through the data points that dete

The SVM algorithm is widely used in machine learning for nonlinear classification tasks. However, when dealing with high-dimensional data, kernel functions are used to transform the data high-dimensional space to achieve better separation. This application of kernel functions depends on the choice of kernel function, such as linear kernels, polynomial kernels, RBF kernels, or sigmoid kernels, depending on the specific use case.



Types of SVM classifiers

Linear SVMs

Linear SVMs are used with linearly separable data. They do not undergo any transformations to separate the two classes. Your privacy choices are respected by the fact that support vectors form the primary boundary.

Winston from MIT uses the analogy of “[fitting](#) this quadratic optimization problem. Mathem represented as:

$$wx + b = 0$$

where w is the weight vector, x is the input ve

There are two approaches to calculating the r classes, which are hard-margin classification hard-margin SVMs, the data points will be per vectors, or “off the street” to continue with Pi represented with the formula,

$$(wx_i + b) y_i \geq a,$$

and then the margin is maximized, which is re the margin projected onto w .

Soft-margin classification is more flexible, allowing use of slack variables (` ξ `). The hyperparameter γ narrows the margin for minimal misclassification, allowing for more misclassified data³.

Nonlinear SVMs

Much of the data in real-world scenarios are nonlinear, so nonlinear SVMs come into play. In order to make the model more complex, various preprocessing methods are applied to the training data to increase the dimensionality of the feature space. That said, higher dimensional spaces can lead to increased complexity by increasing the risk of overfitting, which is computationally taxing. The “kernel trick” helps to reduce computation more efficiently, and it does this by mapping the data into an equivalent kernel function⁴.

There are a number of different kernel types! Some popular kernel functions include:

- [Polynomial kernel](#)
- [Radial basis function](#), kernel (also known as RBF)

- Sigmoid kernel

Support vector regression

Support vector regression (SVR) is an extension of SVMs to regression problems (i.e. the outcome is continuous). Similar to SVMs, the goal of SVR is to find a hyperplane with the maximum margin between data points, but instead of classification, it performs regression prediction.

SVR differs from linear regression in that you are looking to understand between the independent variables, not just the dependent variable. This is because linear regression only provides a linear understanding of the relationships between variables, while SVR can capture non-linear relationships on their own.

Industry newsletter

The latest AI trends, by experts

Get curated insights on the most important to our weekly Think newsletter. See the [IBN](#)

johndoe@yourdomain.com



How SVMs work

In this section, we will discuss the process of other supervised learning algorithms and its applications.

Building a SVM classifier

Your privacy choices

Split your data

As with other machine learning models, start testing set. As an aside, this assumes that you **analyze** on your data. While this is technically good practice before using any machine learning model, it's also important to understand of any missing data or outliers.

Generate and evaluate the model

Import an SVM module from the library of your choice. Fit the training samples on the classifier and predict the test samples. Then, evaluate the accuracy by comparing accuracy of the test set to the predicted values. You can also evaluate other evaluation metrics, like f1-score, precision, or recall.

Hyperparameter tuning

Hyperparameters can be tuned to improve the performance of the model. Common hyperparameters can be found using grid search or random search. You can iterate through different kernel, regularization, and optimization algorithm combination.

SVMs vs. other supervised learning models

Different machine learning classifiers can be compared and contrasted. You can test out and evaluate different models to understand their strengths and weaknesses. Said that, it can be helpful to understand the strengths and weaknesses of each model for your specific application for your use case.

SVMs vs naive bayes

Both Naive Bayes and SVM classifiers are commonly used for classification tasks. SVMs tend to perform better than Naive Bayes in most cases. That said, SVMs have to tune for different hyperparameters, which makes them computationally expensive.

Your privacy choices
SVMs vs logistic regression

SVMs typically perform better with high-dimensional image and text data, compared to logistic regression. They are less prone to overfitting and easier to interpret. That said, t

SVMs vs decision trees

SVMs perform better with high-dimensional data compared to decision trees. That said, decision trees are particularly effective with smaller datasets, and they are more interpretable.

SVM vs. neural networks

Similar to other model comparisons, SVMs are faster and less prone to overfitting, but neural networks are more scalable.

Mixture of Experts | 21 November, epis



Decoding AI: Weekly News Roundup

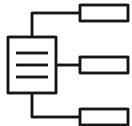
Join our world-class panel of engineers and more as they cut through the AI news and insights.

[Watch all episodes of Mixture of Experts](#)

Your privacy choices

Applications of SVMs

While SVMs can be applied for a number of tasks, here are some applications of SVMs across industries.



Text classification

SVMs are commonly used in natural language processing tasks such as sentiment analysis, spam filtering, and topic modeling. They lend themselves to these data as they are often high-dimensional data.

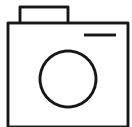


Image classification

SVMs are applied in image classification tasks such as object detection and image retrieval. It can also be useful in security domains, classifying an image as one that has been tampered with.



Bioinformatics

SVMs are also used for protein classification, gene expression analysis, and disease diagnosis. SVMs are often applied in [cancer research](#) because they can detect subtle trends in complex datasets.



Geographic information system (GIS)

SVMs can analyze layered geophysical structures underground, filtering out the 'noise' from electromagnetic data. They have also helped to predict the seismic liquefaction potential of soil, which is relevant to field of civil engineering.

Your privacy choices

Ebook

Unlock the power of generative AI + ML

Learn how to confidently incorporate generative AI and machine learning into your business.

[Read the ebook](#)



Resources

Report

IBM is named a Leader in Data Science & Machine Learning

Learn why IBM has been recognized
Your privacy choices
as a Leader in the 2025 Gartner®

Magic Quadrant™ for Data Science and Machine Learning Platforms.

[Read the report](#)



AI models

Explore IBM Granite

IBM Granite is our family of open, performant and trusted AI models, tailored for business and optimized to scale your AI applications.

Explore language, code, time series and guardrail options.

[Meet Granite](#)



Report

AI in Action 2024

We surveyed 2,000 organizations about their AI initiatives to discover what's working, what's not and how you can get ahead.

[Read the report](#)



Your privacy choices

Explainer

Supervised learning models

Explore supervised learning approaches such as support vector machines and probabilistic classifiers.

Training

Hands-on with generative AI

Learn fundamental concepts and build your skills with hands-on labs, courses, guided projects, trials and more.

[Learn generative AI](#) →

Your privacy choices

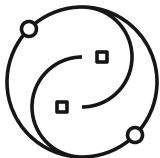
Ebook

~~How to choose the right~~

1 / 2



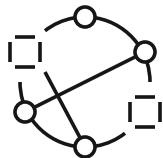
Related solutions



IBM® watsonx Orchestrate™

Easily design scalable AI assistants and agents, automate repetitive tasks and simplify complex processes with IBM® watsonx Orchestrate™.

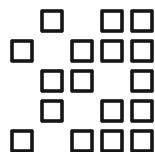
[Explore watsonx Orchestrate →](#)



Artificial intelligence solutions

Put AI to work in your business with IBM's industry-leading AI expertise and portfolio of solutions at your side.

[Explore AI solutions →](#)



Artificial intelligence consulting and services

IBM Consulting AI services help reimagine how businesses work with AI for transformation.

[Explore AI services →](#)

Take the next step

Your privacy choices

Whether you choose to customize pre-built apps and skills or build and deploy custom agentic services using an AI studio, the IBM watsonx platform has you covered.

[Explore watsonx Orchestrate](#) →

[Explore watsonx.ai](#) →

Your privacy choices